

# SpikingConvNeXt: TTFS Conversion of ConvNeXt with Minimum Residual Merging and Learnable Delays for Sparse Inference

<Author 1>, <Author 2>, ...

<Affiliation(s)>

<Corresponding author email>

Prepared for submission to <Target Journal Name> (Q1).

## Abstract

Time-to-first-spike (TTFS) coding is an attractive paradigm for energy-efficient spiking neural networks (SNNs) because each neuron emits at most one spike, and more salient information is encoded by earlier spikes. Recent approximation-free TTFS theory and training methods can match deep ReLU accuracy while operating with high spiking sparsity [2]. In this work, we propose SpikingConvNeXt: a conversion of ConvNeXt [1] into a TTFS-SNN that (i) removes activation and normalization layers, (ii) replaces pointwise MLP transformations with an analytic spike-time mapping, and (iii) introduces a time-domain residual merge based on elementwise minimum rather than summation. The minimum residual merge is motivated by the interpretation of early spikes as high-saliency evidence: in the time domain, adding residual paths delays spikes and can suppress critical early information, while a minimum operator preserves the earliest arriving evidence. To control sparsity, we further introduce learnable per-output delays in the TTFS mapping and optionally constrain pointwise weights to be non-negative to bias neurons toward late or absent spikes. On CIFAR-10, SpikingConvNeXt-Tiny reaches 95.96% top-1 accuracy with 33.3% activation sparsity and achieves up to ~38% activation sparsity with a modest accuracy trade-off (94.43%). Results for CIFAR-100 and MNIST are left as placeholders to be completed.

**Keywords:** Spiking neural networks; time-to-first-spike; TTFS; ConvNeXt; sparsity; residual connections; neuromorphic computing

## Contributions

A TTFS-based spiking conversion of ConvNeXt blocks that removes GELU and normalization layers while reusing ConvNeXt weights. A time-domain residual merge operator based on elementwise minimum that improves information flow for spike-time representations. Learnable per-output delays (and optional non-negative pointwise weights) to directly steer activation sparsity in deep TTFS networks. A practical PyTorch implementation and a reproducible evaluation protocol on CIFAR-10/100 and MNIST (placeholders included).

## 1. Introduction and Related Work

Spiking neural networks (SNNs) offer an event-driven alternative to conventional artificial neural networks by communicating with discrete spikes. In time-to-first-spike (TTFS) coding, each neuron emits at most one spike within a fixed time window, and information is represented by the spike time; earlier spikes typically indicate stronger evidence. This representation supports low-latency inference because decisions can be formed from early activity, while inactive neurons remain silent and are encoded as late spike times (often clamped to  $t_{\max}$ ) [2].

Modern vision backbones, however, are designed for dense real-valued activations. ConvNeXt is a strong example: it achieves high accuracy with a clean ConvNet design augmented by effective training practices from transformer-era models [1]. Its blocks rely on value-domain nonlinearities (e.g., GELU), normalization layers, and additive residual merges. These operations are not naturally defined over spike times. Applying them directly to spike-time tensors undermines TTFS semantics, and additive merges can obscure the interpretation that earlier spikes represent more salient evidence.

This work introduces SpikingConvNeXt, a TTFS formulation of ConvNeXt that operates explicitly in the spike-time domain. We remove GELU and normalization layers and reuse the ConvNeXt depthwise and pointwise weights while replacing the value-domain MLP pathway with an analytic TTFS mapping adopted from prior approximation-free TTFS formulations [2]. Depthwise convolution is applied to spike-time maps, followed by two pointwise projections implemented by closed-form spike-time updates. Input images are encoded as spike times using a linear intensity-to-time mapping within a fixed window  $[t_{\min}, t_{\max}]$ .

A key architectural change is the residual merge. Standard ConvNeXt combines the residual and transformed paths by addition. In TTFS, the earliest spike often carries the strongest evidence, so we replace residual addition with an elementwise minimum over spike times. This minimum-based fusion preserves the earliest evidence across the skip and transformed paths while keeping the representation within TTFS semantics. Figure 3 provides an overview of additive versus minimum residual fusion and Figure 4 illustrates its effect with a numerical example.

We further introduce learnable per-output delays in both pointwise TTFS projections ( $D_{\text{mid}}$  and  $D_{\text{out}}$ ). These delays shift spike times toward later values and increase the frequency of clamping to  $t_{\max}$ , which raises activation sparsity. We also study an optional non-negative constraint on pointwise weights to encourage later or absent spikes under TTFS dynamics. Together, these modifications bridge a high-performing ConvNet backbone [1] with an analytic TTFS computation model [2], and provide simple mechanisms to preserve TTFS semantics and steer sparsity in deep networks.

## 2. Method

### 2.1. ConvNeXt backbone and modifications

ConvNeXt blocks consist of a depthwise convolution, followed by a pointwise MLP implemented as two  $1 \times 1$  linear transformations with a nonlinearity (GELU) in between, and a residual connection [1]. In our implementation, we remove normalization layers and replace GELU with an analytic spike-time mapping

inspired by the TTFS conversion framework [2]. We reuse the ConvNeXt weights for the depthwise and pointwise layers and wrap each block with a spike-time computation module.

**Figure 1.**




Figure 1. End-to-end pipeline of SpikingConvNeXt: input images are encoded into TTFS spike times, processed by a SpikingConvNeXt backbone, pooled, converted to scores, and mapped to logits.

**Figure 2.**

#### Spiking ConvNeXt block (proposed)




Figure 2. Proposed spiking variant of a ConvNeXt block. The depthwise convolution operates on spike-time maps, while the pointwise transforms are implemented by an analytic TTFS mapping with learnable delays. The residual merge uses an elementwise minimum operator.

## 2.2. TTFS encoding of images

Given an input image tensor  $I \in [0,1]^{C \times H \times W}$ , we encode each pixel intensity into a spike time within a fixed time window  $[t_{\min}, t_{\max}]$  using a linear TTFS encoding:

$$t = t_{\min} + (1 - I) \cdot (t_{\max} - t_{\min}).$$

Higher intensity values map to earlier spike times (closer to  $t_{\min}$ ), while low intensities map to later times (closer to  $t_{\max}$ ). In our implementation, a spike time clamped to  $t_{\max}$  is treated as “no spike” for sparsity accounting.

### 2.3. Analytic spike-time mapping for pointwise layers

To replace the ANN pointwise MLP (Linear  $\rightarrow$  GELU  $\rightarrow$  Linear) with a TTFS operation, we compute output spike times directly from input spike times. For a vector of input spike times  $t_j$ , a weight matrix  $W$ , and a per-output delay  $D_i$ , we compute:

- (1)  $\text{threshold} = t_{\max} - t_{\min} - D_i$
- (2)  $t_i = (t_j - t_{\min}) \cdot W + \text{threshold} + t_{\min}$
- (3)  $t_i = \min(t_i, t_{\max})$

This operation corresponds to the PyTorch function `call_spiking_torch` used in our implementation. Delays  $D_i$  are learned and constrained to be non-negative and bounded within the time window.

---

**Algorithm 2. Analytic TTFS mapping used in pointwise layers.**

---

**Algorithm 2 CallSpiking: Analytic TTFS Spike-Time Mapping**

*Input:*  $t_j \in \mathbb{R}^{B \times C_{\text{in}}}$  (input spike times),  $W \in \mathbb{R}^{C_{\text{in}} \times C_{\text{out}}}$  (weights),  $D \in \mathbb{R}^{C_{\text{out}}}$  (delay),  $t_{\min}, t_{\max}$

*Output:*  $t_i \in \mathbb{R}^{B \times C_{\text{out}}}$  (output spike times)

- 1:  $\theta \leftarrow (t_{\max} - t_{\min}) - D$  // threshold term (Eq. 18 form)
  - 2:  $\Delta \leftarrow t_j - t_{\min}$
  - 3:  $t_i \leftarrow (\Delta \cdot W) + \theta + t_{\min}$
  - 4:  $t_i \leftarrow \text{where}(t_i < t_{\max}, t_i, t_{\max})$  // clamp late/non-spikes to  $t_{\max}$
  - 5: return  $t_i$
- 

### 2.4. SpikingConvNeXt block

Each ConvNeXt block is converted into a SpikingBlock that reuses the original depthwise and pointwise weights. Given an input spike-time tensor  $t_j \in \mathbb{R}^{N \times C \times H \times W}$ , we apply the depthwise convolution directly on the spike-time map as a practical spatial mixing heuristic. We then flatten spatial positions and apply the analytic TTFS mapping twice, corresponding to the two pointwise linear layers. We introduce two learnable delay vectors,  $D_{\text{mid}}$  and  $D_{\text{out}}$ , for the hidden and output pointwise transforms.

---

**Algorithm 1. Spiking ConvNeXt block forward pass (TTFS mapping with min-residual).**

---

**Algorithm 1 Algorithm of Spiking ConvNeXt Block**

*Input:*  $t_j \in \mathbb{R}^{N \times C \times H \times W}$  (input spike times),  $t_{\min}, t_{\max}$ ;  $DWConv(\cdot)$ ;  $W1, W2$ ;  $D_{\text{mid}}, D_{\text{out}}$ ;  
 $DropPath(\cdot)$ ;  $force\_positive$   
*Output:*  $t_{\text{out}} \in \mathbb{R}^{N \times C \times H \times W}$  (output spike times)

```
1:  $x \leftarrow DWConv(t_j)$ 
2:  $(N, C, H, W) \leftarrow \text{shape}(x)$ 
3:  $x_{\text{flat}} \leftarrow \text{reshape}(\text{permute}(x, N, H, W, C), (N \cdot H \cdot W, C))$ 
4: if  $force\_positive$  then
5:    $W1 \leftarrow \text{ReLU}(W1)$ ;  $W2 \leftarrow \text{ReLU}(W2)$  // enforce non-negative pointwise weights
6: end if
7:  $D_{\text{mid}} \leftarrow \text{clamp}(\text{ReLU}(D_{\text{mid}}), 0, 0.9 \cdot (t_{\max} - t_{\min}))$  // learnable per-output delays
8:  $t_{\text{mid}} \leftarrow \text{CallSpiking}(x_{\text{flat}}, W1^T, D_{\text{mid}}, t_{\min}, t_{\max})$ 
9:  $D_{\text{out}} \leftarrow \text{clamp}(\text{ReLU}(D_{\text{out}}), 0, 0.9 \cdot (t_{\max} - t_{\min}))$ 
10:  $t_{\text{blk}} \leftarrow \text{CallSpiking}(t_{\text{mid}}, W2^T, D_{\text{out}}, t_{\min}, t_{\max})$ 
11:  $t_{\text{blk}} \leftarrow \text{permute}(\text{reshape}(t_{\text{blk}}, (N, H, W, C)), N, C, H, W)$ 
12:  $t_{\text{out}} \leftarrow \min(t_j, DropPath(t_{\text{blk}}))$  // time-domain residual fusion
13: return  $t_{\text{out}}$ 
```

---

**Innovations in this block.** Compared with the original ConvNeXt block, we (i) remove normalization and GELU and replace the pointwise MLP with an analytic TTFS mapping (Algorithm 2), (ii) fuse the residual and block output in the time domain using an elementwise minimum instead of summation so earlier spikes dominate, and (iii) introduce learnable non-negative delays ( $D_{\text{mid}}, D_{\text{out}}$ ) and optionally non-negative pointwise weights to push spike times toward  $t_{\max}$ , increasing silent-neuron rate.


## 2.5. Minimum residual merge in the time domain

In standard ConvNeXt (and most ANNs), the residual merge is additive:  $y = x + f(x)$ . In a TTFS network, signals represent spike times rather than activation magnitudes, and addition does not preserve the semantics of latency. Moreover, adding spike times generally delays the resulting value, potentially suppressing early evidence. We therefore propose an elementwise minimum residual merge:


$$t_{\text{out}} = \min(t_j, f_{\text{TTFS}}(t_j)).$$

Figure 3 contrasts the standard additive residual merge with the proposed minimum operator in the TTFS time domain; Figure 4 provides a numeric illustration.

This choice is consistent with TTFS semantics where earlier spike times indicate more salient evidence [2]. The minimum operator selects the earliest available evidence across the shortcut and transformed paths.



**Figure 3.** Residual fusion in ConvNeXt (addition) versus SpikingConvNeXt in the TTFS spike-time domain (minimum). The minimum operator preserves earliest evidence; late or absent spikes are represented by clamping to  $t_{max}$ .



**Figure 4.** Illustration of residual merging in the time domain. ANN-style summation delays spike times, while the proposed minimum merge preserves the earliest (most salient) spike.

## 2.6. Sparsity mechanisms

We incorporate two mechanisms to steer activation sparsity in deep TTFS networks:

**Learnable delays:**  $D_{mid}$  and  $D_{out}$  are trained to control the effective timing threshold per output neuron. **Optional non-negative weights:** we optionally enforce  $W1$  and  $W2$  to be non-negative (via ReLU on weights) to bias outputs toward late spikes and increase the fraction of neurons clamped at  $t_{max}$ .

## 2.7. Why non-negative pointwise weights can increase sparsity

Let  $\Delta_j = (t_j - t_{min}) \geq 0$  denote the time offset from the window start. From Eq. (2), the pre-clamp output is  $t_i = \Delta_j \cdot W + (t_{max} - t_{min} - D_i) + t_{min}$ . When  $\Delta_j \geq 0$  and  $W \geq 0$ , the product term  $\Delta_j \cdot W$  is non-negative, which shifts  $t_i$  toward larger values (later spikes). More values therefore exceed  $t_{max}$  and are clamped to  $t_{max}$ , which we count as silent (no spike) events. This provides an intuitive mechanism to bias the network toward sparse spiking activity, at the cost of reduced representational flexibility if enforced too strongly.

Activation sparsity is measured as the fraction of neurons that do not fire within the time window, operationalized as spike times equal to  $t_{\max}$  (after clamping). If  $T$  is the set of all spike-time values in a layer for a batch, sparsity is computed as:

$$S = |\{t \in T : t = t_{\max}\}| / |T|.$$

### 3. Experiments

#### 3.1. Datasets

We evaluate on CIFAR-10, CIFAR-100, and MNIST. CIFAR-10 and CIFAR-100 are  $32 \times 32$  natural image datasets with 10 and 100 classes, respectively. MNIST is a  $28 \times 28$  grayscale digit dataset with 10 classes. All inputs are scaled to  $[0,1]$  and encoded into spike times using the TTFS encoder in Section 3.2.

#### 3.2. Model variants

Unless otherwise stated, we use ConvNeXt-Tiny as the backbone configuration (depths = (3, 3, 9, 3), dims = (96, 192, 384, 768)) [1]. We consider the following variants:

Try 1 (baseline conversion): Replace the GELU pointwise nonlinearity with TTFS mapping; keep residual merge behavior unchanged. Try 2 (minimum residual): Use the proposed elementwise minimum residual merge; delay terms fixed to  $D = 0$ . Try 3 (learnable delays): Use minimum residual merge and enable learning of  $D_{\text{mid}}$  and  $D_{\text{out}}$  to increase activation sparsity.

#### 3.3. Training and evaluation protocol

Training details are provided as placeholders to be completed:

- Initialization: <trained from scratch / initialized from pretrained ConvNeXt weights>
- Optimizer: <e.g., SGD / AdamW>, learning rate: <...>, weight decay: <...>
- Augmentation: <random crop, flip, cutout, mixup, ...>
- Epochs: <...>, batch size: <...>, hardware: <GPU/TPU details>

During evaluation, we use a fixed time window  $t_{\min} = 0.0$  and  $t_{\max} = 1.0$  and report top-1 accuracy and activation sparsity.

### 4. Results

#### 4.1. CIFAR-10 results

Table 1 summarizes CIFAR-10 test results from our experiments. The proposed minimum residual merge (Try 2) improves accuracy over the baseline spiking conversion (Try 1). Enabling learnable delays (Try 3) increases activation sparsity to approximately 38% with a modest accuracy reduction.

Table 1. Test accuracy and activation sparsity (CIFAR-10 filled; CIFAR-100 and MNIST left as placeholders).

| Method                         | CIFAR-10<br>Acc. (%) | CIFAR-10<br>Sparsity (%) | CIFAR-100<br>Acc. (%) | CIFAR-100<br>Sparsity (%) | MNIST Acc.<br>(%) | MNIST<br>Sparsity (%) |
|--------------------------------|----------------------|--------------------------|-----------------------|---------------------------|-------------------|-----------------------|
| Nazari et al.<br>[placeholder] | 93.60                | -                        | <TBD>                 | <TBD>                     | <TBD>             | <TBD>                 |

|                                       |       |       |       |       |       |       |
|---------------------------------------|-------|-------|-------|-------|-------|-------|
| ref]                                  |       |       |       |       |       |       |
| Simonyan (ANN) [placeholder ref]      | 93.56 | 0 (?) | <TBD> | <TBD> | <TBD> | <TBD> |
| Stanojevic et al. (TTFS baseline) [2] | 93.59 | 38.00 | <TBD> | <TBD> | <TBD> | <TBD> |
| Ours (SNN) Try 1                      | 95.29 | 33.30 | <TBD> | <TBD> | <TBD> | <TBD> |
| Ours (SNN) Try 2 (min residual)       | 95.96 | 33.30 | <TBD> | <TBD> | <TBD> | <TBD> |
| Ours (SNN) Try 3-A (learnable delays) | 94.11 | 38.18 | <TBD> | <TBD> | <TBD> | <TBD> |
| Ours (SNN) Try 3-B (learnable delays) | 94.43 | 38.01 | <TBD> | <TBD> | <TBD> | <TBD> |
| Ours (SNN) Try 3-C (learnable delays) | 94.05 | 38.50 | <TBD> | <TBD> | <TBD> | <TBD> |

**Figure 4.**






Figure 4. Accuracy vs. sparsity trade-off on CIFAR-10. Minimum-residual variants achieve higher accuracy at comparable sparsity, while learnable delays push sparsity upward with a modest accuracy trade-off.

**Figure 5.**




Figure 5. Example stage-wise activation sparsity (Try 3-B snapshot). Early stages exhibit low sparsity compared to later stages, highlighting opportunities to further sparsify early feature extraction.

## 4.2. Discussion

Why does minimum residual help? In TTFS coding, an earlier spike implies a larger effective activation (higher saliency) [2]. In an additive residual, combining spike times with summation produces larger latencies, which can suppress early evidence. The minimum residual merge avoids this artificial delay and acts as a simple deterministic rule to preserve early evidence.

Sparsity control. Learnable delays increase activation sparsity by allowing the model to shift some neurons toward late spikes that are clamped to  $t_{\max}$ , effectively silencing them. The observed stage-wise sparsity suggests that later stages are easier to sparsify than early stages; future work can explore stage-specific delays or additional regularizers to increase sparsity in early layers without harming accuracy.

## 5. Limitations and Future Work

This manuscript is a draft intended for journal submission and contains placeholders that must be completed (training details, CIFAR-100 and MNIST results, statistical significance over multiple runs, and baseline citations).

Methodologically, the depthwise convolution is applied directly to spike-time maps as a practical heuristic. A more principled TTFS-compatible convolution operator may further improve accuracy and sparsity. Additionally, enforcing strictly non-negative pointwise weights may increase sparsity but can limit representational capacity; the best constraint likely depends on the dataset and should be explored with ablations.

For a Q1 submission, we recommend: (i) report mean $\pm$ std over  $\geq 3$  seeds, (ii) provide energy/latency proxies (spikes/neuron, silent ratio, and optionally FLOPs or hardware measurements), (iii) compare against strong modern SNN baselines and ANN ConvNeXt, and (iv) include an ablation over  $t_{\min}/t_{\max}$  and delay initialization.

## 6. Conclusion

We introduced SpikingConvNeXt, a TTFS-based conversion of ConvNeXt [1] that removes activations and normalization, replaces pointwise MLP computations with an analytic spike-time mapping, and proposes a minimum residual merge suited to time-domain representations. The approach achieves strong CIFAR-10 performance (up to 95.96% top-1) while enabling controllable sparsity through learnable delays. These results suggest that modern ConvNet backbones can be adapted to TTFS SNNs with minimal structural changes and simple time-domain operators, paving the way for energy-efficient deployment on neuromorphic and event-driven hardware.

## Acknowledgments

<Acknowledgments text (funding, compute grants, collaborators, etc.)>

## Data Availability

All datasets used are publicly available (CIFAR-10/100 and MNIST). Exact preprocessing and splits: <provide details / repository link>.

## Code Availability

Source code and trained checkpoints will be released at: <GitHub/Zenodo link> upon publication.

## Appendix A: Key Implementation Snippets

This appendix includes short code snippets for clarity. For full source code, see the repository link in Code Availability.

### Listing A1: Analytic TTFS mapping (call\_spiking\_torch).

```
def call_spiking_torch(tj, w, D_i, t_min_prev, t_min, t_max):  
    # threshold (Eq. 18 in the reference TTFS formulation)  
    threshold = t_max - t_min - D_i
```

```

tj = tj.to(W.dtype)
threshold = threshold.to(W.dtype)
delta = tj - t_min
ti = torch.matmul(delta, W) + threshold + t_min
ti = torch.where(ti < t_max, ti, t_max)
return ti

```

### Listing A2: Time-domain residual merge.

```

# ANN residual:
# out = tj + self.drop_path(t_out)

# Proposed TTFS residual merge:
out = torch.minimum(tj, self.drop_path(t_out))

```

## References

- [1] Liu, Z., Mao, H., Wu, C.-Y., Feichtenhofer, C., Darrell, T., & Xie, S. A ConvNet for the 2020s (ConvNeXt). arXiv:2201.03545, 2022.
- [2] Stanojevic, A., Gehring, W., Woźniak, S., Binas, J., Chicca, E., & Payeur, A. High-performance deep spiking neural networks with 0.3 spikes per neuron. Nature Communications, 15:6793, 2024.
- [3] Krizhevsky, A. Learning multiple layers of features from tiny images. Technical report (CIFAR), 2009.
- [4] LeCun, Y., Cortes, C., & Burges, C. J. C. The MNIST database of handwritten digits, 1998.
- [5] <Add additional SNN/TTFS references used in Related Work and baselines.>